

MACH



Classification



MNIST Dataset as a running example

- A set of 70,000 small images of digits handwritten by high school students and employees of the US Census Bureau.
 - ◆ Each image is labeled with the digit it represents.
 - ◆ A.k.a. the « hello world » of Machine Learning.

```
from sklearn.datasets import fetch_openml  
  
mnist = fetch_openml('mnist_784', as_frame=False)
```




- as_frame = False to get the data as Numpy arrays instead of Panda DataFrame.
- 70,000 images, each image has 28x28 (features). Each feature simply represents the pixel intensity (white:0 to black:255)



The classification task

- Learn a model to recognize digits
- Create a test set and set it aside before inspecting the data closely.
- This dataset is **already shuffled**.



```
X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]
```



Binary Classification



Training a Binary classifier

- We simplify the problem and we want to identify only one digit.
 - E.g., the '5'
- The 5-detector is an example of a binary classifier capable of distinguishing between just two classes: 5 and not-5.

```
y_train_5 = (y_train == '5')  
y_test_5 = (y_test == '5')
```

- Next step: pick a classifier and train it.
 - Stochastic Gradient Descent (SGD)
 - Capable of handling large dataset efficiently
 - Suited for Online learning.

```
from sklearn.linear_model import SGDClassifier  
  
sgd_clf = SGDClassifier(random_state=42)  
sgd_clf.fit(X_train, y_train_5)
```



Performance measures

- Evaluating a classifier is often trickier than evaluating a regressor
- Many performance measures



Measuring accuracy using cross-validation

- K-fold cross-validation:
 - splitting the training set into k folds
 - Then, training the model k times, holding out a different fold each time for the evaluation.
 - 3-fold cross-validation on MNIST:
 - [0.95035, 0.96035, 0.9604]
 - Above 95% accuracy (ratio of correct predictions) on all cv folds !!!
 - Amazing ! Do you believe it ?
 - Look at a dummy classifier (just classify every image in the most frequent class:
 - [0.90965, 0.90965, 0.90965]
- ➔ Accuracy is generally not the preferred performance measure for classifier, especially when dealing with **skewed** datasets (unbalanced).



Confusion Matrices

- Counting the number of times instances class A are classified as class B, for all A/B pairs.

| | | Predicted class | |
|--------------|-------|-----------------|------|
| | | Non-5 | 5 |
| Actual class | Non-5 | 53892 | 687 |
| | 5 | 1891 | 3530 |

| | | Predicted class | |
|--------------|----------|-----------------|----------|
| | | Negative | Positive |
| Actual class | Negative | TN | FP |
| | Positive | FN | TP |

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$



Precision / Recall: intuitions

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$




F measure to combine precision and recall

- It is often convenient to combine precision and recall into a single metric
- F_1 score is the harmonic mean (more weight to low values) of precision and recall
 - $F_1 = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{TP}{TP + \frac{FN + FP}{2}}$
- This measure favors classifiers with similar precision and recall.
- Increasing precision reduces recall and vice versa: aka precision/ recall trade-off



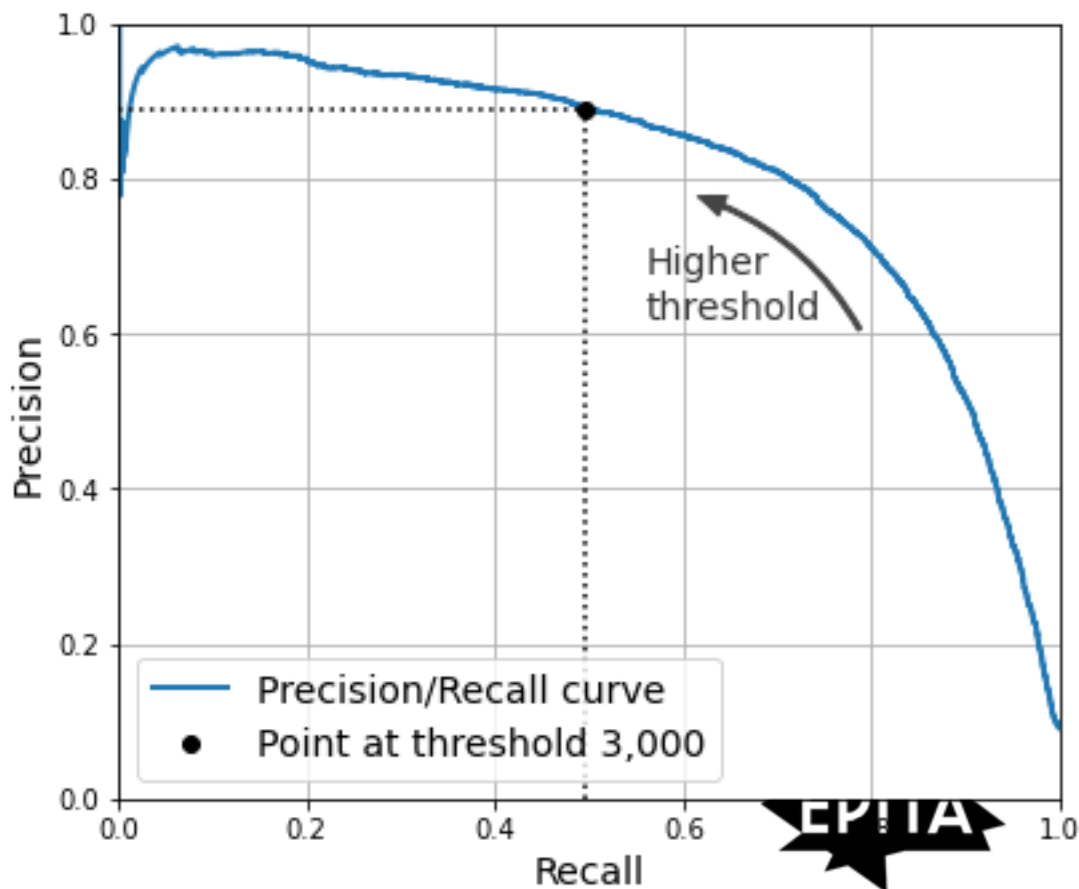
The precision/recall trade-off

- SGD classifier (and not only) computes a score based on a **decision function**.
 - If this score is **greater** than a **threshold**, it assigns the instance to the **positive class**; otherwise it assigns it to the **negative class**.



Precision / recall curve

- This plot can be used to set the threshold
- E.g., one wants a 90% precisions
 - Search for the lowest threshold that gives at least 90% precision.
- Easy to create a classifier with desired precision
 - What about the recall?




The ROC Curve

- The **receiver operating characteristic** (ROC) is a common tool used with binary classifier.
 - Very similar to precision/recall curve
 - The ROC curve plots the **TP rate** (recall) vs the **FP rate** (also called the *fall-out*)
 - FPR= ratio of negative instances that are incorrectly classified as positive.
 - $FPR = 1 - TNR$
 - TNR: ratio of negative instances that are correctly classified as negative.
 - TNR: also called specificity
- ROC curve plots sensitivity (recall) versus 1-specificity.



- Again a trade-off
- One way to compare classifier is to measure the area under the curve (AUC).
- A perfect classifier will have a ROC AUC equal to 1.
- A purely random:
 - ROC AUC=0.5




ROC curve or PR curve ?

- They are similar so: **how to decide which one to use ?**
- Prefer PR curve whenever
 - **the positive class is rare**
 - Or you **care more about the false positives than the false negatives**.
- Otherwise, use the ROC curve (and ROC AUC score).
- Example:
 - Considering the previous ROC curve you may think that the classifier is really good but this is mostly because there are few positives (5s) compared to the negatives (non-5s). In contrast, PR curve makes it clear that the classifier has room for improvement.



Considering another classifier.

- Random Forest
- ROC AUC : 0.99 (vs 0.96 for SGD)



Multiclass Classification

To distinguish between more than two classes



To distinguish between more than two classes

- Aka: *multinomial* classifiers
- Some classifiers are able to handle multiple classes **natively** (e.g., Logistic reg., Random Forest, GaussianNB)
- Others are **strictly binary classifiers** (e.g., SGD, SVC)
- ➔ There are **various strategies** to perform multiclass classification with **multiple binary classifiers**.
 - ➔ *One-Versus-All* (OVA/OVR)
 - ➔ *One-Versus-One* (OVO)



One-versus-the-rest / one-versus-all (OVR/OVA)

- Create a system that can classify the instances into k classes by training k binary classifiers.
 - One classifier for each class.
 - To classify a new instance:
 - take the decision for each classifier
 - Select the class whose classifier outputs the **highest** score.
- MNIST: train 10 binary classifiers (one for each digit)
 - 0-detector, 1-detector, ..., 9-detector



One-versus-One

- Train a binary classifier for every pair of labels.
 - One to distinguish 0s and 1s, another to distinguish 0s and 2s ...
 - Need to train 45 binary classifiers
 - If N classes $\rightarrow (N \times (N-1) / 2)$ classifiers.
- To classify an image:
 - Run the image through all the classifiers
 - See which class wins the most duels.

+ : each classifier only needs to be trained on the part of the training set containing the two classes it must distinguish.

- Some algorithms (e.g., SVM) scale poorly with the size of the training set \rightarrow OvO is preferred because it is faster to train many classifiers on small training set than few classifiers on large training sets.




- Scikit-Learn detect when a binary classification algorithm is used for a multiclass classification task.
- To force Scikit-Learn to use OvO or OVR:
 - OneVsOneClassifier or OneVsRestClassifier classes
 - Simply create an instance and pass a classifier to its constructor.
 - `ovr_clf = OneVsRestClassifier(SVC(random_state=42))`
 - `ovr_clf.fit(X_train[:2000], y[train:2000])`
 - `over_clr.predict([some_digits])`

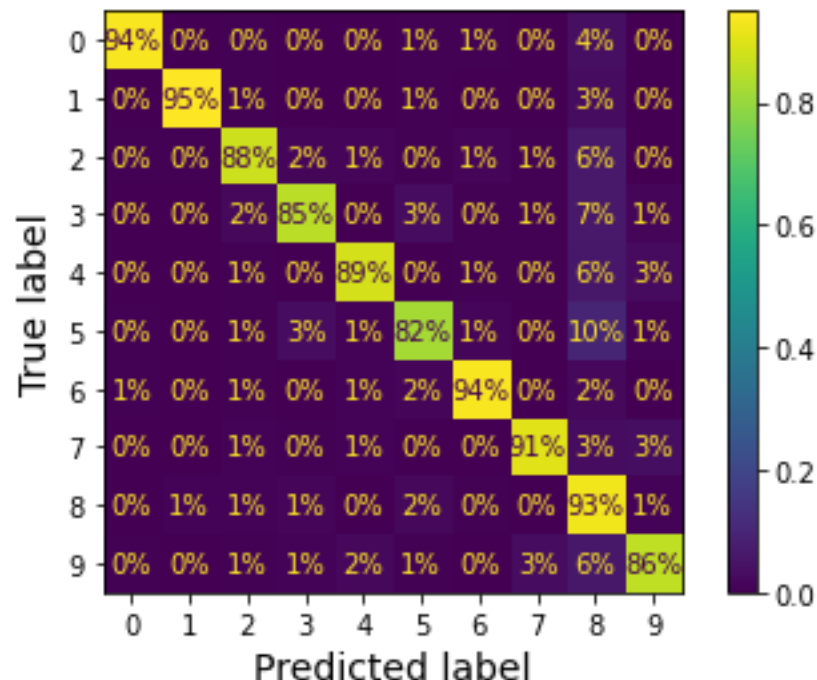


Error Analysis

```
from sklearn.metrics import ConfusionMatrixDisplay  
  
y_train_pred = cross_val_predict(sgd_clf, X_train_scaled, y_train, cv=3)  
plt.rc('font', size=9) # extra code - make the text smaller  
ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred)  
plt.show()
```



```
plt.rc('font', size=10) # extra code
ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred,
                                       normalize="true", values_format=".0%")
plt.show()
```


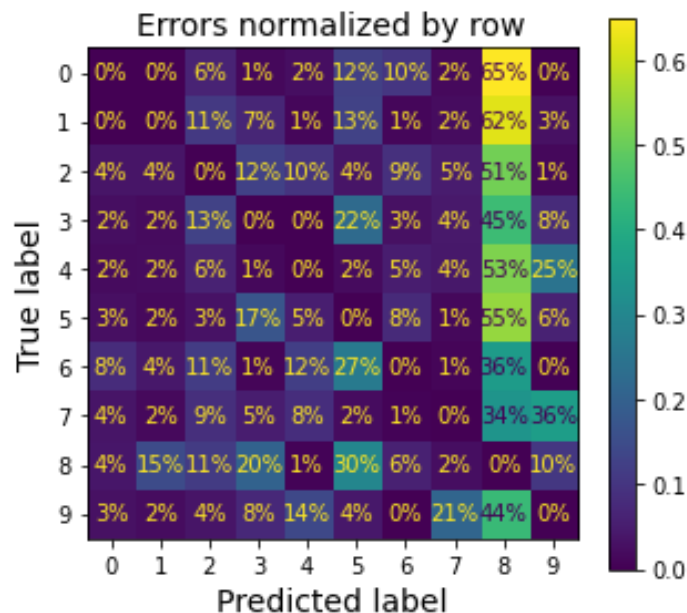


82% of '5's are correctly classified

10% of instances classified as '5's are actually '8's

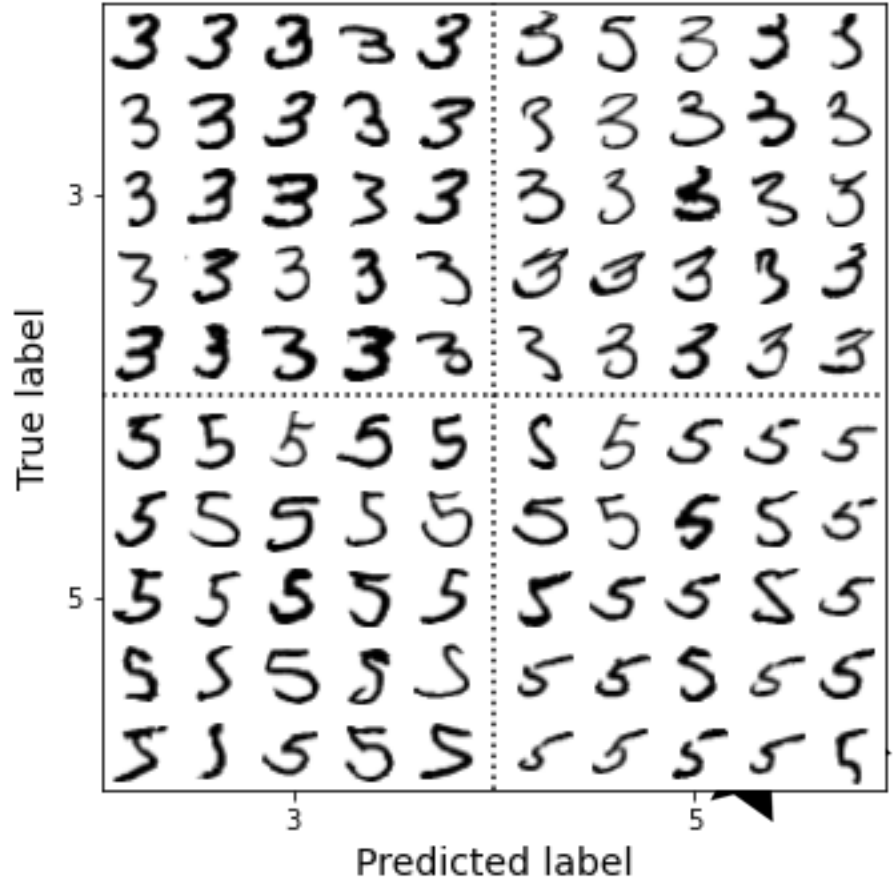


- (0,8) left means 65% of errors the model made on '0's were misclassified as '8's.
- (9,7) right means 56% of misclassified '7's are actually '9's.



Go further

- The classifier is quite sensitive to:
 - Image shifting and rotation.
- ➔ Heavy and conscientious preprocessing
- ➔ Or **data augmentation**.



Multilabel Classification



- Until now, each instance has been assigned to just one class.
 - In some cases, you may want your classifier to output multiple classes for each instance.
 - Face recognition: several people in the same picture.
 - News: may have several topics (e.g., diplomacy, sport, politics, business).
- ➔ A system that outputs **multiple binary tags** is called a **multilabel classification system**.



Ex: Large and odd digit ?

```
import numpy as np
from sklearn.neighbors import KNeighborsClassifier

y_train_large = (y_train >= '7')
y_train_odd = (y_train.astype('int8') % 2 == 1)
y_multilabel = np.c_[y_train_large, y_train_odd]

knn_clf = KNeighborsClassifier()
knn_clf.fit(X_train, y_multilabel)
```

- To go further: have a look at ChainClassifier to capture dependency between labels.






Multiclass Classification



Multi-output Classification

- Multioutput-multiclass Classification or just Multioutput Classification
- A generalization of multilabel classification where each label can be multiclass (i.e., can have more than two possible values).
- Example: image denoising



Summary



Conclusion

- Now, you know:
 - How to select good metrics for classification tasks,
 - Pick the appropriate precision/recall trade-off,
 - Compare classifiers,
 - Build good classification systems on a variety of tasks.
- Next level:
 - Learn how all these machine learning models actually work.



The end

- Exercices
- One mini-project next time